

COSC2406/2407 Database Systems

Tree Index Structures

Assignment Project Exam Help

Xiangmin (Emily) Zhou

<https://powcoder.com>

RMIT University

Room: 14.11.04

Online Consultation via Collaborate Ultra (no appointment required): 10:20-11.20am

Thursdays

Email : xiangmin.zhou@rmit.edu.au

Add WeChat powcoder

Lecture 5

References: Ramakrishnan & Gehrke Chapter 10

Garcia-Molina et al. Chapter 13

Elmasri & Navathe Chapter 5

Diagrams courtesy Ramakrishnan & Gehrke

Assignment Project Exam Help

In this lecture, we will discuss tree index structures.

Specifically, we will discuss:

- The Indexed Sequential Access Method (ISAM)
- Dynamic B+-trees

<https://powcoder.com>

Add WeChat powcoder

- As for any index, 3 alternatives for index data entries k^* :

- ① Data record with search key value k
- ② $(k, \text{rid of data record with search key value } k)$
- ③ $(k, \text{list of rids of data records with search key } k)$

– Choice is orthogonal to the indexing technique.

- Tree-structured indexing techniques support both range searches and equality searches
- ISAM (Indexed Sequential Access Method) is a tree-based static structure
- B+-tree is a dynamic structure that adjusts with insertions and deletions

Motivation for Tree Structures: Range Searches

- *"Find all students older than 22."*

If data is in sorted file, we can binary search to find first such student, then scan to find other matches


- The cost of binary search can be high
- Simple solution: create an "index" file (see Section 8.2)



- Now we can binary search on the (smaller) index file!

Motivation for Tree Structures: Range Searches ...

Index file may still be quite large. Apply idea repeatedly.



Leaf pages contain data entries.

Assignment Project Exam Help

Consider a heap file, with 2 records per page:

<https://powcoder.com>

Data
Pages

2*	5*	27*	55*	3*	37*	10*	46*	51*	33*	40*	97*
----	----	-----	-----	----	-----	-----	-----	-----	-----	-----	-----

Add WeChat powcoder

Assignment Project Exam Help

Sort file on search key:

<https://powcoder.com>

Data
Pages

1*	5*	20*	27*	33*	37*	40*	46*	51*	55*	63*	97*
----	----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----

Add WeChat powcoder

Build an index on file: **Assignment Project Exam Help**

<https://powcoder.com>



Build a sparse index on file:

Assignment Project Exam Help

<https://powcoder.com>



Build an index on the index:


Assignment Project Exam Help

<https://powcoder.com>

Index

Data
Pages

Add WeChat powcoder




The tree we have constructed is a so-called ISAM (Indexed Sequential Access Method) structure. (See Section 10.2)

Each index page has *index data entries* of the form (key, pointer).

Each index page contains one pointer more than the number of keys.

A key serves as a *separator* between the contents of the pages pointed to by the pointers to its left and right.

index entry



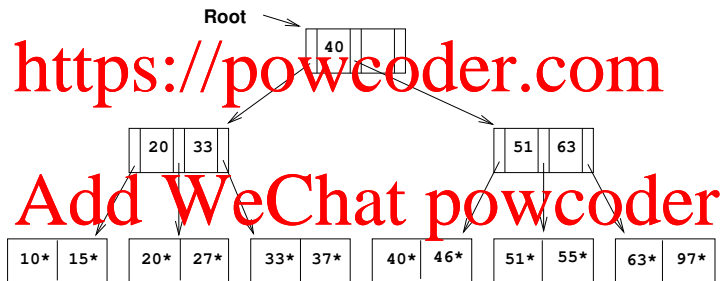
- File creation: Leaf (data) pages allocated sequentially, sorted by search key; then index pages allocated, then space for overflow pages.

- Index entries: (search key value, page id); these “direct” the search for data entries, which are in leaf pages.
- Search: Start at the root; use key comparisons to go to a leaf. Cost: $\log_F N$, where F is *fan-out* of the index page and N is number of leaf pages.


- Insert: Find the leaf page where the data entry belongs to and put it there. Create overflow pages if necessary.
- Delete: Find and remove the entry from the leaf; if this empties an overflow page, de-allocate the page.

Example ISAM Tree

Each node can hold 2 entries no need for 'next-leaf-page' pointers
(because of the sequential allocation of leaf pages).



After Inserting 23*, 48*, 41*, 42* ...




Advantage: Less locking problems in ISAM than in other structures, since no index page is changed.

... Then Deleting 42*, 51*, 97*

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Note that 51* appears in index levels, but not in a leaf!

Assignment Project Exam Help

Once an ISAM file is created, inserts and deletes affect only the leaf page content.

As a result, long overflow chains can develop. This affects the time required to retrieve a record, since the chains have to be searched as well.

To reduce this problem, the tree can be created with some free space (20%, say) for future insertions.

Otherwise, the only way to get rid of overflow chains requires a reorganisation of the whole file structure.

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

The fact that only leaf pages are adjusted gives a big advantage for *concurrent* access:

When a page is accessed, it is usually *locked* to ensure that it is not concurrently modified by another user. This can result in long queues of users waiting to access a page. Such a situation can cause significant performance issues, especially if the locked node is near the root of a tree.

Since ISAM index pages are never modified, they do not need to be locked.


B+-tree: The Most Widely Used Index

The B+-tree is the most widely used index structure (see Section 10.3). It has the following characteristics:

- Internal nodes direct the search, index data entries are in the leaf
- The tree is kept *height-balanced*
- Insert and delete at $\log_F N$ cost, where F is the fan-out and N the number of leaves
- Minimum 50% occupancy of nodes (except for the root).
- Each node contains $d \leq m \leq 2d$ entries, where the parameter d is called the *order* of the tree.
- Supports equality and range-searches efficiently
- Leaf pages are organised as a double linked list for fast traversal and reorganisation

Assignment Project Exam Help

<https://powcoder.com>



Add WeChat powcoder

Non-leaf nodes contain m index entries, with $m + 1$ pointers to children (like SAM structure).

Leaf nodes contain data entries, either:

- actual data records (alternative 1); the B+-tree is an *integrated index*, where the file contains the index structure as well as the data
- pointers to data records elsewhere on disk (alternatives 2 & 3); the B+-tree is an *index file*, distinct from the file which contains the records


If the indexed field is of fixed-length, so is the index entry. Otherwise, the index entries are variable-length.

Assignment Project Exam Help

- A search begins at the root, and key comparisons direct the search to a leaf (as in ISAM)
- (Try a search for 5*, 15*, and all data entries $\geq 24^*$)

<https://powcoder.com>

Add WeChat powcoder




- Find correct leaf L
- Put data entry into L
 - If L has enough space, done!
 - Else, must *split* L (into L and a new node $L2$)
 - Redistribute entries evenly, copy up middle key.
 - Insert index entry pointing to $L2$ into parent of L .
- This can happen recursively.
 - To split index node, redistribute entries evenly, but *push-up* middle key. (Contrast to leaf splits!)
- Splits “grow” tree; root split increases height.
 - Tree growth: gets wider or one level taller at top.


Inserting 8* into Example B+-tree

- Observe how minimum occupancy level is guaranteed in both leaf and index pages splits.

- Note difference between *copy-up* and *push-up*!




<https://powcoder.com>



Example B+-tree After Inserting of 8*

Assignment Project Exam Help



- Notice that root was split, leading to increase in height.
- In this example, we can avoid split by re-distributing entries; however, this is usually not done in practice.

Deleting a Data Entry from a B+-Tree

- Start at root; find the leaf L where the entry belongs.

- Remove the entry.

- If L is at least half full, done!

- if L has only $d - 1$ entries,

- Try to re-distribute, borrowing from *sibling* (adjacent node with same parent as L) - If re-distributing fails, merge L and sibling.

- If merge occurred, must delete entry (pointing to L or sibling) from parent of L .

- Merge could propagate to root, decreasing height.

Refer to Section 10.6 for an example.

Garcia-Molina et al. illustrates a B-tree example in Section 13.3.6.

Assignment Project Exam Help



Add WeChat powcoder

Deletion of 20* is solved by re-distribution from sibling nodes. Note that the new splitting (middle) key is copied up.

Deleting 24*


Leaf nodes must be merged:

Assignment Project Exam Help

<https://powcoder.com>

Then, index nodes need to be merged and entry 17 is pulled down:

Add WeChat powcoder



Assignment Project Exam Help

- Typical order (minimum node size): 100. Typical fill-factor: 67%
 - Average fan-out = 133
(increased by *key compression*)
- Typical Capacities:
 - Height 4: $133^4 = 312,900,700$ records
 - Height 3: $133^3 = 2,352,637$ records
- Can often hold top levels in the buffer manager's pool:
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 Mbytes

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Read more indexes in MongoDB (including use of B+-trees)

<https://docs.mongodb.com/manual/indexes>

<https://powcoder.com>

Read more about how B+-trees are implemented in Derby:

http://db.apache.org/derby/papers/btree_package.html

Add WeChat powcoder

Assignment Project Exam Help

We have discussed tree index structures in this lecture.

Specifically, we have discussed Indexed Sequential Access Method (ISAM) and dynamic B+-trees.

<https://powcoder.com>
Add WeChat powcoder